

Problem Statement

Scaler is an online tech-versity offering intensive computer science & Data Science courses through live classes delivered by tech leaders and subject matter experts. The meticulously structured program enhances the skills of software professionals by offering a modern curriculum with exposure to the latest technologies. It is a product by InterviewBit

```
In [1]: import numpy as np
import pandas as pd
import re
import matplotlib.pyplot as plt
import seaborn
from sklearn.preprocessing import MinMaxScaler
```

```
In [2]: df = pd.read_csv('scaler_clustering.csv')
```

```
In [3]: df.head()
```

Out[3]:

| | Unnamed: 0 | company_hash | email_hash | orgyear | ctc | job_position | ctc_updated_year |
|---|------------|--------------------------|---|---------|---------|--------------------|------------------|
| 0 | 0 | atrgrxnnnt xzaxv | 6de0a4417d18ab14334c3f43397fc13b30c35149d70c05... | 2016.0 | 1100000 | Other | 2020.0 |
| 1 | 1 | qtrxvzwt xzegwgbbrxbxnta | b0aaf1ac138b53cb6e039ba2c3d6604a250d02d5145c10... | 2018.0 | 449999 | FullStack Engineer | 2019.0 |
| 2 | 2 | ojzwnvwnxw vx | 4860c670bcd48fb96c02a4b0ae3608ae6fdd98176112e9... | 2015.0 | 2000000 | Backend Engineer | 2020.0 |
| 3 | 3 | ngpgutaxv | effdede7a2e7c2af664c8a31d9346385016128d66bbc58... | 2017.0 | 700000 | Backend Engineer | 2019.0 |
| 4 | 4 | qxen sqghu | 6ff54e709262f55cb999a1c1db8436cb2055d8f79ab520... | 2017.0 | 1400000 | FullStack Engineer | 2019.0 |

```
In [4]: df.shape
```

Out[4]: (205843, 7)

```
In [5]: df=df.drop(columns='Unnamed: 0')
df=df.drop(columns='email_hash')
```

```
In [6]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205843 entries, 0 to 205842
Data columns (total 5 columns):
 #   Column                Non-Null Count  Dtype  
---  -
 0   company_hash          205799 non-null object  
 1   orgyear                205757 non-null float64  
 2   ctc                   205843 non-null int64  
 3   job_position          153281 non-null object  
 4   ctc_updated_year      205843 non-null float64  
dtypes: float64(2), int64(1), object(2)
memory usage: 7.9+ MB
```

```
In [7]: df.describe()
```

```
Out[7]:
```


| | orgyear | ctc | ctc_updated_year |
|--------------|---------------|--------------|------------------|
| count | 205757.000000 | 2.058430e+05 | 205843.000000 |
| mean | 2014.882750 | 2.271685e+06 | 2019.628231 |
| std | 63.571115 | 1.180091e+07 | 1.325104 |
| min | 0.000000 | 2.000000e+00 | 2015.000000 |
| 25% | 2013.000000 | 5.300000e+05 | 2019.000000 |
| 50% | 2016.000000 | 9.500000e+05 | 2020.000000 |
| 75% | 2018.000000 | 1.700000e+06 | 2021.000000 |
| max | 20165.000000 | 1.000150e+09 | 2021.000000 |

```
In [8]: df.isnull().sum()/len(df)*100
```

```
Out[8]: company_hash      0.021376  
orgyear      0.041779  
ctc          0.000000  
job_position  25.534995  
ctc_updated_year  0.000000  
dtype: float64
```


```
In [9]: fig = plt.figure(figsize = (15, 5))  
df['orgyear'].value_counts().head(25).plot(kind='bar')
```

```
Out[9]: <AxesSubplot:>
```




```
In [10]: fig = plt.figure(figsize = (10, 5))  
df['ctc_updated_year'].value_counts().plot(kind='bar')
```

Out[10]: <AxesSubplot:>




```
In [11]: fig = plt.figure(figsize = (10, 5))  
df['ctc'].value_counts().head(25).plot(kind='kde')
```

Out[11]: <AxesSubplot:ylabel='Density'>



```
In [12]: fig = plt.figure(figsize = (15, 5))  
df['job_position'].value_counts().head(25).plot(kind='bar')
```

Out[12]: <AxesSubplot:>



Data Pre-processing

```
In [13]: _data_nums=df.orgyear  
         #keeping only the numerical columns
```

```
In [14]: _data_nums=_data_nums.reset_index()
```

```
In [15]: _data_nums.drop(columns='index',inplace=True)
```

```
In [16]: columns=_data_nums.columns
```

```
In [17]: from sklearn.impute import KNNImputer  
         imputer = KNNImputer(n_neighbors=5, weights='uniform', metric='nan_euclidean',)  
         imputer.fit(_data_nums)  
         # transform the dataset  
         _data_new = imputer.transform(_data_nums)
```

```
In [18]: _data_new=pd.DataFrame(_data_new)
```

```
In [19]: _data_new.columns=columns
```

```
In [20]: _data_new.isnull().sum()
```

```
Out[20]: orgyear      0  
         dtype: int64
```

Getting the remaining columns back

```
In [21]: remaining_columns=list(set(df.columns).difference(set(columns)))  
         data=pd.concat([_data_new, df[remaining_columns]],axis=1)  
         data=pd.DataFrame(data)
```

```
In [22]: data.dropna(inplace=True)
```

Regex for cleaning company names

```
In [23]: try:
          for i in range(len(data)):
              data['company_hash'][i]=re.sub('[^A-Za-z0-9 ]+', '', data['company_hash'][i])
          except Exception:
              pass
```

/var/folders/2m/svsbyfss4h53t29lk1vcnvf40000gn/T/ipykernel_7376/1537211381.py:3: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
data['company_hash'][i]=re.sub('[^A-Za-z0-9 ]+', '', data['company_hash'][i])
```

New Feature Creation

```
In [24]: data['years_of_Experience']=2023-data['orgyear']

data = data[(data['years_of_Experience'] >= 0) & (data['years_of_Experience'] <= 100)]
```



```
In [25]: data
```

```
Out[25]:
```

| | orgyear | job_position | company_hash | ctc_updated_year | ctc | years_of_Experience |
|--------|---------|--------------------|---------------------------|------------------|---------|---------------------|
| 0 | 2016.0 | Other | atrgxnnt xzaxv | 2020.0 | 1100000 | 7.0 |
| 1 | 2018.0 | FullStack Engineer | qtrxvzwt xzegwgbb rxbxnta | 2019.0 | 449999 | 5.0 |
| 2 | 2015.0 | Backend Engineer | ojzwnvwnxw vx | 2020.0 | 2000000 | 8.0 |
| 3 | 2017.0 | Backend Engineer | ngpgutaxv | 2019.0 | 700000 | 6.0 |
| 4 | 2017.0 | FullStack Engineer | qxen sqghu | 2019.0 | 1400000 | 6.0 |
| ... | ... | ... | ... | ... | ... | ... |
| 205324 | 2016.0 | FullStack Engineer | wos xzntqzvnxgzvr | 2021.0 | 1500000 | 7.0 |
| 205326 | 2019.0 | FullStack Engineer | xzegojo | 2021.0 | 1200000 | 4.0 |
| 205327 | 2015.0 | Data Scientist | wgbuzgcv wgzngqwn | 2021.0 | 1000000 | 8.0 |
| 205328 | 2019.0 | Data Scientist | ahzzyhbmj | 2021.0 | 1100000 | 4.0 |
| 205329 | 2017.0 | Frontend Engineer | ertdnqvat ojointbo rbn | 2021.0 | 1100000 | 6.0 |

153178 rows × 6 columns

Standardization & Encoding

```
In [26]: X = data
```

```
y = data['job_position']
```

```
In [27]: from sklearn.preprocessing import LabelEncoder

le = LabelEncoder()

X['job_position'] = le.fit_transform(X['job_position'])

y = le.transform(y)
```

/var/folders/2m/svsbyfss4h53t29lklvcnvf40000gn/T/ipykernel_7376/199861872.py:5: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)
X['job_position'] = le.fit_transform(X['job_position'])

```
In [28]: X = X
y = X['company_hash']
```

```
In [29]: from sklearn.preprocessing import LabelEncoder

le = LabelEncoder()

X['company_hash'] = le.fit_transform(X['company_hash'])

y = le.transform(y)
```

/var/folders/2m/svsbyfss4h53t29lklvcnvf40000gn/T/ipykernel_7376/1388851734.py:6: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)
X['company_hash'] = le.fit_transform(X['company_hash'])

```
In [30]: data=X
```

```
In [31]: data
```

```
Out[31]:
```

| | orgyear | job_position | company_hash | ctc_updated_year | ctc | years_of_Experience |
|---------------|---------|--------------|--------------|------------------|---------|---------------------|
| 0 | 2016.0 | 458 | 870 | 2020.0 | 1100000 | 7.0 |
| 1 | 2018.0 | 292 | 18095 | 2019.0 | 449999 | 5.0 |
| 2 | 2015.0 | 140 | 14223 | 2020.0 | 2000000 | 8.0 |
| 3 | 2017.0 | 140 | 11066 | 2019.0 | 700000 | 6.0 |
| 4 | 2017.0 | 292 | 18560 | 2019.0 | 1400000 | 6.0 |
| ... | ... | ... | ... | ... | ... | ... |
| 205324 | 2016.0 | 292 | 28211 | 2021.0 | 1500000 | 7.0 |
| 205326 | 2019.0 | 292 | 31034 | 2021.0 | 1200000 | 4.0 |
| 205327 | 2015.0 | 208 | 27490 | 2021.0 | 1000000 | 8.0 |
| 205328 | 2019.0 | 208 | 489 | 2021.0 | 1100000 | 4.0 |
| 205329 | 2017.0 | 287 | 5388 | 2021.0 | 1100000 | 6.0 |

153178 rows × 6 columns

```
In [32]: cols = X.columns
```

```
In [33]: from sklearn.preprocessing import MinMaxScaler
```

```
ms = MinMaxScaler()
```

```
X = ms.fit_transform(X)
```

```
In [34]: X = pd.DataFrame(X, columns=[cols])
```

```
In [35]: X.head()
```

```
Out[35]:
```

| | orgyear | job_position | company_hash | ctc_updated_year | ctc | years_of_Experience |
|---|----------|--------------|--------------|------------------|---------|---------------------|
| 0 | 0.867925 | 0.450787 | 0.025444 | 0.833333 | 0.00550 | 0.132075 |
| 1 | 0.905660 | 0.287402 | 0.529202 | 0.666667 | 0.00225 | 0.094340 |
| 2 | 0.849057 | 0.137795 | 0.415962 | 0.833333 | 0.01000 | 0.150943 |
| 3 | 0.886792 | 0.137795 | 0.323633 | 0.666667 | 0.00350 | 0.113208 |
| 4 | 0.886792 | 0.287402 | 0.542801 | 0.666667 | 0.00700 | 0.113208 |

Manual Clustering

```
In [36]: sf_company_hash = data.groupby('company_hash').agg({'ctc' : ['mean', 'median', 'max', 'min', 'count']}).reset_in
```

```
In [37]: sf_company_hash
```

```
Out[37]:
```

| | company_hash | ctc | | | | |
|--------------|--------------|-----------|-----------|---------|---------|-------|
| | | mean | median | max | min | count |
| 0 | 0 | 100000.0 | 100000.0 | 100000 | 100000 | 1 |
| 1 | 1 | 300000.0 | 300000.0 | 300000 | 300000 | 1 |
| 2 | 2 | 550000.0 | 550000.0 | 830000 | 270000 | 2 |
| 3 | 3 | 1100000.0 | 1100000.0 | 1100000 | 1100000 | 1 |
| 4 | 4 | 250000.0 | 250000.0 | 250000 | 250000 | 1 |
| ... | ... | ... | ... | ... | ... | ... |
| 34189 | 34189 | 2400000.0 | 2400000.0 | 2400000 | 2400000 | 1 |
| 34190 | 34190 | 940000.0 | 940000.0 | 940000 | 940000 | 1 |
| 34191 | 34191 | 1370000.0 | 1370000.0 | 1370000 | 1370000 | 1 |
| 34192 | 34192 | 600000.0 | 600000.0 | 600000 | 600000 | 1 |
| 34193 | 34193 | 720000.0 | 720000.0 | 720000 | 720000 | 1 |

34194 rows × 6 columns

```
In [38]: sf_job_position = data.groupby('job_position').agg({'ctc' : ['mean', 'median', 'max', 'min', 'count']}).reset_in
```

```
In [39]: sf_job_position
```

```
Out[39]:
```

| | job_position | ctc | | | | | |
|------|--------------|-----------|-----------|---------|---------|-------|--|
| | | mean | median | max | min | count | |
| 0 | 0 | 1200000.0 | 1200000.0 | 1200000 | 1200000 | 1 | |
| 1 | 1 | 700000.0 | 700000.0 | 700000 | 700000 | 1 | |
| 2 | 2 | 600000.0 | 600000.0 | 600000 | 600000 | 1 | |
| 3 | 3 | 470000.0 | 470000.0 | 470000 | 470000 | 1 | |
| 4 | 4 | 420000.0 | 420000.0 | 420000 | 420000 | 1 | |
| ... | ... | ... | ... | ... | ... | ... | |
| 1012 | 1012 | 1715000.0 | 1715000.0 | 2400000 | 1030000 | 2 | |
| 1013 | 1013 | 2000000.0 | 2000000.0 | 2000000 | 2000000 | 1 | |
| 1014 | 1014 | 500000.0 | 500000.0 | 500000 | 500000 | 1 | |
| 1015 | 1015 | 610000.0 | 610000.0 | 610000 | 610000 | 1 | |
| 1016 | 1016 | 82000.0 | 82000.0 | 82000 | 82000 | 1 | |

1017 rows × 6 columns

```
In [40]: _Experience = data.groupby('years_of_Experience').agg({'ctc' : ['mean', 'median', 'max', 'min', 'count']}).reset
```

```
In [41]: sf_years_of_Experience
```

```
Out[41]:
```

| | years_of_Experience | ctc | | | | |
|---|---------------------|--------------|----------|-----------|-------|-------|
| | | mean | median | max | min | count |
| 0 | 0.00000 | 1.429270e+07 | 740000.0 | 200000000 | 10000 | 179 |
| 1 | 1.00000 | 7.209064e+06 | 800000.0 | 200000000 | 2000 | 605 |
| 2 | 2.00000 | 4.247075e+06 | 650000.0 | 200000000 | 3400 | 2377 |
| 3 | 3.00000 | 2.268745e+06 | 670000.0 | 200000000 | 24 | 6955 |
| 4 | 4.00000 | 1.633791e+06 | 700000.0 | 200000000 | 1000 | 13868 |
| 5 | 5.00000 | 1.815167e+06 | 750000.0 | 200000000 | 500 | 19004 |
| 6 | 6.00000 | 2.083714e+06 | 800000.0 | 200000000 | 1000 | 17928 |
| 7 | 7.00000 | 2.332430e+06 | 869999.0 | 200000000 | 25 | 17839 |
| 8 | 8.00000 | 2.432603e+06 | 950000.0 | 200000000 | 1000 | 16179 |
| 9 | 8.11725 | 2.732952e+06 | 905000.0 | 100000000 | 8000 | 62 |

```
In [42]: data
```

```
Out[42]:
```

| | orgyear | job_position | company_hash | ctc_updated_year | ctc | years_of_Experience |
|---------------|---------|--------------|--------------|------------------|---------|---------------------|
| 0 | 2016.0 | 458 | 870 | 2020.0 | 1100000 | 7.0 |
| 1 | 2018.0 | 292 | 18095 | 2019.0 | 449999 | 5.0 |
| 2 | 2015.0 | 140 | 14223 | 2020.0 | 2000000 | 8.0 |
| 3 | 2017.0 | 140 | 11066 | 2019.0 | 700000 | 6.0 |
| 4 | 2017.0 | 292 | 18560 | 2019.0 | 1400000 | 6.0 |
| ... | ... | ... | ... | ... | ... | ... |
| 205324 | 2016.0 | 292 | 28211 | 2021.0 | 1500000 | 7.0 |
| 205326 | 2019.0 | 292 | 31034 | 2021.0 | 1200000 | 4.0 |
| 205327 | 2015.0 | 208 | 27490 | 2021.0 | 1000000 | 8.0 |
| 205328 | 2019.0 | 208 | 489 | 2021.0 | 1100000 | 4.0 |
| 205329 | 2017.0 | 287 | 5388 | 2021.0 | 1100000 | 6.0 |

153178 rows × 6 columns

Unsupervised learning


```
In [43]: data=pd.DataFrame(data, columns=['orgyear','job_position','company_hash','ctc','ctc_updated_year','years_of_Exp
```

K- means clustering


```
In [44]: from sklearn.cluster import KMeans

Inter = []
for i in range(1,20):
    model = KMeans(n_clusters = i)
    model.fit(X)
    Inter.append(model.inertia_)

# plotting the Elbow
plt.figure(figsize = (14, 8))
plt.plot(np.arange(1,20), Inter)
plt.xlabel('Clusters')
plt.ylabel('SSE')
plt.show()
```



Hierarchical Clustering

```
In [45]: from sklearn.decomposition import PCA
pca = PCA(n_components = 6)
X_principal = pca.fit_transform(X)
X_principal = pd.DataFrame(X_principal)
X_principal.columns = ['P1', 'P2', 'P3', 'P4', 'P5', 'P6']

X_principal.head(2)
```

Out[45]:

| | P1 | P2 | P3 | P4 | P5 | P6 |
|---|-----------|-----------|-----------|-----------|-----------|---------------|
| 0 | 0.480169 | -0.151548 | 0.142753 | -0.022936 | -0.007883 | 9.909772e-14 |
| 1 | -0.003887 | 0.071486 | -0.003421 | -0.098304 | -0.006352 | -6.177225e-16 |

```
In [ ]: import scipy.cluster.hierarchy as shc
plt.figure(figsize =(6, 6))
plt.title('Visualising the data')
Dendrogram = shc.dendrogram((shc.linkage(X_principal, method ='ward')))
```

Actionable Insights & Recommendations

In []:

In []:

In []: